

## Runtime Registries

AIOSv1 system hosts several registries for keeping track of the running blocks, clusters and vDAGs.

These are the following registries in AIOSv1:

1. **Clusters registry** - Clusters registry stores the information of all the clusters that are currently on-boarded into the network.
2. **Blocks registry** - Blocks registry stores the information of all the blocks that are currently running across all the clusters in the network.
3. **vDAGs registry** - vDAGs registry stores all the vDAGs that are currently created in the network across multiple blocks.
4. **vDAG controller registry** - vDAG controller registry stores all the vDAG controllers that are created to serve vDAG inference requests.

### Clusters registry:

Clusters registry contains the information about the on-boarded clusters in the network.

Here is the schema of the cluster entry in clusters registry:

```
const clusterSchema = new Schema({
  // Unique ID for the cluster
  id: { type: String, required: true, unique: true },

  // Optional region or network ID where the cluster is deployed
  regionId: { type: String, required: false },

  status: { type: String, required: true },

  // Aggregated and per-node information for all nodes in the cluster
  nodes: {
    // Total number of nodes in the cluster
    count: { type: Number, required: true },
    // Detailed info for each individual node
    nodeData: [{
      // Unique ID for the node
      id: { type: String, required: true },
      // GPU details for the node
      gpus: {
        // Number of GPUs in the node
        count: { type: Number, required: true },
        // Total GPU memory in MB
        memory: { type: Number, required: true },
      },
    ]
  }
})
```

```

// List of GPU models with individual memory sizes
gpus: [{
    modelName: { type: String, required: true }, // GPU model name
    memory: { type: Number, required: true } // Memory per GPU in MB
}],
// Optional GPU features (e.g., CUDA versions)
features: [String],
// List of distinct GPU model names
modelName: [String]
},
// Virtual CPU details
vcpus: {
    count: { type: Number, required: true } // Number of vCPUs in the node
},
// Total memory in MB
memory: { type: Number, required: true },
// Total swap space in MB
swap: { type: Number, required: true },
// Storage info per node
storage: {
    disks: { type: Number, required: true }, // Number of disks
    size: { type: Number, required: true } // Total storage size in MB
},
// Network interface stats per node
network: {
    interfaces: { type: Number, required: true }, // Number of network interfaces
    txBandwidth: { type: Number, required: true }, // Transmit bandwidth (MBps)
    rxBandwidth: { type: Number, required: true } // Receive bandwidth (MBps)
}
}],
},

// Total GPU stats across all nodes
gpus: {
    count: { type: Number, required: true }, // Total number of GPUs in the cluster
    memory: { type: Number, required: true } // Total GPU memory in MB
},

// Total vCPU count across the cluster
vcpus: {
    count: { type: Number, required: true }
},

// Total memory across the cluster in MB
memory: { type: Number, required: true },

```

```

// Total swap space across the cluster in MB
swap: { type: Number, required: true },

// Aggregated storage details for the cluster
storage: {
  disks: { type: Number, required: true }, // Total number of disks
  size: { type: Number, required: true } // Total storage size in MB
},

// Aggregated network configuration
network: {
  interfaces: { type: Number, required: true }, // Total number of interfaces
  txBandwidth: { type: Number, required: true }, // Total TX bandwidth
  rxBandwidth: { type: Number, required: true } // Total RX bandwidth
},

// Configuration used by the cluster controller
config: {
  type: new Schema({
    policyExecutorId: { type: String, required: false, default: "" }, // Option
    policyExecutionMode: { type: String, required: false, default: "local" }, // Exe
    customPolicySystem: { type: Schema.Types.Mixed, required: false }, // Any c
    publicHostname: { type: String, required: true }, // Public
    useGateway: { type: Boolean, required: false, default: true }, // Wheth
    actionsPolicyMap: { type: Schema.Types.Mixed, required: false }, // Mappin
    // URLs to internal/external services in the cluster
    urlMap: {
      controllerService: { type: String, required: true }, // URL for control
      metricsService: { type: String, required: true }, // URL for metrics
      blocksQuery: { type: String, required: true }, // URL for querying
      publicGateway: { type: String, required: true }, // Public-facing g
      parameterUpdater: { type: String, required: true } // URL for model/c
    }
  }),
  required: true
},

// List of user-defined tags or labels
tags: { type: [String], required: true },

// Human-readable metadata about the cluster
clusterMetadata: {
  type: new Schema({
    name: { type: String, required: true }, // Friendly name of th
    description: { type: String, required: true }, // Purpose or use-case
    owner: { type: String, required: true }, // Who owns or manages

```

```

        email: { type: String, required: false }, // Optional contact email
        countries: { type: [String], required: false }, // Countries associated with the cluster
        miscContactInfo: { type: Schema.Types.Mixed, required: false }, // Additional contact information
        additionalInfo: { type: Schema.Types.Mixed, required: false } // Any extra metadata
    }},
    required: true
},

// Reputation score or reliability indicator for the cluster (not yet used anywhere in the system)
reputation: { type: Number, required: false }
});

```

Example:

```

{
  "id": "cluster-west-vision-001",
  "regionId": "us-west-2",
  "status": "live",
  "nodes": {
    "count": 2,
    "nodeData": [
      {
        "id": "node-1",
        "gpus": {
          "count": 2,
          "memory": 32768,
          "gpus": [
            { "modelName": "NVIDIA A100", "memory": 16384 },
            { "modelName": "NVIDIA A100", "memory": 16384 }
          ],
          "features": ["fp16", "tensor_cores"],
          "modelName": ["NVIDIA A100"]
        },
        "vcpus": { "count": 32 },
        "memory": 131072,
        "swap": 8192,
        "storage": {
          "disks": 2,
          "size": 1048576
        },
        "network": {
          "interfaces": 2,
          "txBandwidth": 10000,
          "rxBandwidth": 10000
        }
      }
    ],
    "reputation": 0.95
  }
}

```

```

        "id": "node-2",
        "gpus": {
            "count": 1,
            "memory": 16384,
            "gpus": [
                { "modelName": "NVIDIA V100", "memory": 16384 }
            ],
            "features": ["fp16"],
            "modelName": ["NVIDIA V100"]
        },
        "vcpus": { "count": 16 },
        "memory": 65536,
        "swap": 4096,
        "storage": {
            "disks": 1,
            "size": 524288
        },
        "network": {
            "interfaces": 1,
            "txBandwidth": 5000,
            "rxBandwidth": 5000
        }
    }
]
},
"gpus": {
    "count": 3,
    "memory": 49152
},
"vcpus": {
    "count": 48
},
"memory": 196608,
"swap": 12288,
"storage": {
    "disks": 3,
    "size": 1572864
},
"network": {
    "interfaces": 3,
    "txBandwidth": 15000,
    "rxBandwidth": 15000
},
"config": {
    "policyExecutorId": "policy-exec-007",
    "policyExecutionMode": "local",

```

```

    "customPolicySystem": {
      "name": "AdvancedPolicyRunner",
      "version": "2.1.0"
    },
    "publicHostname": "cluster-west-vision-001.company.net",
    "useGateway": true,
    "actionsPolicyMap": {
      "onScaleUp": "evaluate-gpu-availability",
      "onFailure": "notify-admin"
    },

    // these fields are populated by the system:
    "urlMap": {
      "controllerService": "http://cluster-west-vision-001.company.net:32000/controller",
      "metricsService": "http://cluster-west-vision-001.company.net:32000/metrics",
      "blocksQuery": "http://cluster-west-vision-001.company.net:32000/blocks",
      "publicGateway": "http://cluster-west-vision-001.company.net:32000",
      "parameterUpdater": "http://cluster-west-vision-001.company.net:32000/mgmt"
    }
  },
  "tags": ["gpu", "production", "ml", "vision", "us-west"],
  "clusterMetadata": {
    "name": "Sample cluster",
    "description": "Dedicated to serving large-scale computer vision models in production.",
    "owner": "AI Infrastructure Team",
    "email": "ai-infra@company.net",
    "countries": ["USA", "Canada"],
    "miscContactInfo": {
      "pagerDuty": "https://sample-website/ai-clusters",
      "slack": "#ml-infra"
    },
    "additionalInfo": {

    }
  },
  "reputation": 94
}

```

### Creating a cluster:

For creating the cluster, refer to the documentation of Parser.

### Cluster registry APIs:

**Endpoint:** /clusters/:id

**Method:** GET

**Description:**

Fetches a single cluster document by its unique id.

**Example curl Command:**

```
curl -X GET http://<server-url>/clusters/cluster-west-vision-001
```

---

**Endpoint:** /clusters/:id**Method:** PUT**Description:**

Updates a cluster document by its id using the payload provided in the request body. The body should use MongoDB-style update syntax.

**Example curl Command:**

```
curl -X PUT http://<server-url>/clusters/cluster-west-vision-001 \
-H "Content-Type: application/json" \
-d '{
  "$set": {
    "tags": ["gpu", "updated"],
    "reputation": 97
  }
}'
```

---

**Endpoint:** /clusters/:id**Method:** DELETE**Description:**

Deletes the cluster document with the specified id.

**Example curl Command:**

```
curl -X DELETE http://<server-url>/clusters/cluster-west-vision-001
```

---

**Endpoint:** /clusters/query**Method:** POST**Description:**

Queries cluster documents using a MongoDB-style filter provided in the request body. Supports standard MongoDB operators such as \$eq, \$gt, \$in, etc.

**Example curl Command:**

```
curl -X POST http://<server-url>/clusters/query \
-H "Content-Type: application/json" \
-d '{
  "gpus.count": { "$gte": 2 },
  "clusterMetadata.countries": { "$in": ["USA"] }
}'
```

---

## Blocks registry:

Blocks registry stores the information of all the blocks that are currently running across all the clusters in the network.

Here is the schema of the block:

```
const BlockSchema = new mongoose.Schema({
  // Unique identifier for the block
  id: { type: String, required: true, unique: true },

  // the component URI the block is running - taken from component registry
  componentUri: { type: String },

  // The component data of the block - copied from component registry
  component: { type: mongoose.Schema.Types.Mixed },

  // same as componentUri + block-id
  blockUri: { type: String },

  // Human-readable or structured metadata about the block - copied from component
  blockMetadata: { type: mongoose.Schema.Types.Mixed },

  // Policy configuration or rules tied to the block
  policies: { type: mongoose.Schema.Types.Mixed },

  // The cluster data of the block, copied as it is from the cluster registry
  cluster: { type: mongoose.Schema.Types.Mixed },

  // Data used to initialize the block during deployment/startup
  blockInitData: { type: mongoose.Schema.Types.Mixed },

  // Initialization settings (env vars, args, flags, etc.)
  initSettings: { type: mongoose.Schema.Types.Mixed },

  // Parameters required to configure the block's runtime behavior
  parameters: { type: mongoose.Schema.Types.Mixed },

  // Minimum number of instances this block should maintain
  minInstances: { type: Number, required: false },

  // Maximum number of instances allowed for scaling
  maxInstances: { type: Number, required: false },

  // Input interface specification (protocol - follows a template - copied from component,
```



```

    inputProtocol: { type: mongoose.Schema.Types.Mixed },

    // Output interface specification (protocol - copied from component)
    outputProtocol: { type: mongoose.Schema.Types.Mixed }
  });

```

Example:

```

{
  "id": "block-object-detector-001",
  "componentUri": "",
  "component": {},
  "blockUri": "",
  "blockMetadata": {},
  "policies": {
    "resourceAllocator": {
      "policyRuleURI": "policies.resource_allocator.standard:latest",
      "parameters": {},
      "settings": {}
    },
    "loadBalancer": {
      "policyRuleURI": "policies.load_balancer.gateway.load_balancer_sep2:v0.0.1-beta",
      "parameters": {},
      "settings": {}
    },
    "loadBalancerMapper": {
      "policyRuleURI": "policies.load_balancer.mapper.loadbalancer_mapper_oct1:v1.2.0",
      "parameters": {},
      "settings": {}
    },
    "assignment": {
      "policyRuleURI": "policies.assignment.default_strategy:v1.0.3",
      "parameters": {},
      "settings": {}
    },
    "stabilityChecker": {
      "policyRuleURI": "policies.health.stability_checker:v0.1.0",
      "parameters": {},
      "settings": {}
    },
    "autoscaler": {
      "policyRuleURI": "policies.autoscaler.basic_auto_scaler:v0.3.5",
      "parameters": {},
      "settings": {}
    },
    "accessRulesPolicy": {
      "policyRuleURI": "policies.access.control.access_rules_policy:v2.0.0",

```

```

        "parameters": {},
        "settings": {}
    },
    "cluster": {},
    "blockInitData": {},
    "initSettings": {},
    "parameters": {},
    "minInstances": 1,
    "maxInstances": 5,
    "inputProtocol": {},
    "outputProtocol": {}
}

```

### Creating a block:

For creating the block, refer to the documentation of Parser.

### Block registry APIs:

**Endpoint:** /blocks

**Method:** GET

**Description:**

Fetches all block documents in the database.

**Example curl Command:**

```
curl -X GET http://<server-url>/blocks
```

---

**Endpoint:** /blocks/:id

**Method:** GET

**Description:**

Fetches a single block document by its unique id.

**Example curl Command:**

```
curl -X GET http://<server-url>/blocks/block-object-detector-001
```

---

**Endpoint:** /blocks/:id

**Method:** PUT

**Description:**

Updates a block document by its id using MongoDB-style update syntax in the request body.

**Example curl Command:**

```
curl -X PUT http://<server-url>/blocks/block-object-detector-001 \
-H "Content-Type: application/json" \
-d '{
  "$set": {
    "blockMetadata.description": "Updated description for object detection block",
    "minInstances": 2
  }
}'
```

---

**Endpoint:** /blocks/:id

**Method:** DELETE

**Description:**

Deletes the block document with the specified id.

**Example curl Command:**

```
curl -X DELETE http://<server-url>/blocks/block-object-detector-001
```

---

**Endpoint:** /blocks/query

**Method:** POST

**Description:**

Queries block documents using a MongoDB-style filter provided in the JSON body. Supports standard MongoDB operators such as \$eq, \$gt, \$in, etc. Optional options can be passed for sorting, pagination, etc.

**Example curl Command:**

```
curl -X POST http://<server-url>/blocks/query \
-H "Content-Type: application/json" \
-d '{
  "query": {
    "cluster.reputation": { "$gt": 90 },
    "policies.autoscaler.policyRuleURI": { "$ne": "" }
  },
  "options": {
    "sort": { "id": 1 },
    "limit": 10
  }
}'
```

**Endpoint:** /vdag/:vdagURI

**Method:** PUT

**Description:**

Updates fields in the vDAG document identified by the given vdagURI using MongoDB-style update syntax.

**Example curl Command:**

```
curl -X PUT http://<server-url>/vdag/sample-vdag:1.0-stable \
-H "Content-Type: application/json" \
-d '{
    "$set": {
        "status": "active",
        "metadata.owner": "team-ml"
    }
}'
```

---

**Endpoint:** /vdag/:vdagURI

**Method:** DELETE

**Description:**

Deletes the vDAG document identified by the given vdagURI.

**Example curl Command:**

```
curl -X DELETE http://<server-url>/vdag/sample-vdag:1.0-stable
```

---

**Endpoint:** /vdags

**Method:** POST

**Description:**

Queries multiple vDAG documents using a MongoDB-style filter object.

**Example curl Command:**

```
curl -X POST http://<server-url>/vdags \
-H "Content-Type: application/json" \
-d '{
    "status": "pending",
    "metadata.owner": "team-ml"
}'
```

---

## vDAG controllers registry:

vDAG controller registry stores all the vDAG controllers that are created to serve vDAG inference requests.

Here is the schema of a vDAG controller:

```
from dataclasses import dataclass, field
from typing import Dict, List, Any
```

```
@dataclass
```

```
class vDAGController:
```

```
    # Unique identifier for the vDAG controller instance
```

```

vdag_controller_id: str = ''
# Associated vDAG URI this controller is managing
vdag_uri: str = ''
# Publicly accessible URL for interacting with the controller
public_url: str = ''
# Identifier of the cluster where the controller is deployed
cluster_id: str = ''
# Arbitrary metadata for storing additional information
metadata: Dict[str, Any] = field(default_factory=dict)
# Configuration parameters used by the controller
config: Dict[str, Any] = field(default_factory=dict)
# Tags used for search and discovery of the controller
search_tags: List[str] = field(default_factory=list)

```

### Creating a vDAG controller:

For creating the vDAG controller, refer to the documentation of Parser.

### vDAG controllers registry APIs:

**Endpoint:** /vdag-controller/:controller\_id

**Method:** GET

**Description:**

Fetches the vDAG Controller document identified by the given controller\_id.

**Example curl Command:**

```
curl -X GET http://<server-url>/vdag-controller/controller-123
```

---

**Endpoint:** /vdag-controller/:controller\_id

**Method:** PUT

**Description:**

Updates fields in the vDAG Controller document identified by the given controller\_id using MongoDB-style update syntax.

**Example curl Command:**

```

curl -X PUT http://<server-url>/vdag-controller/controller-123 \
-H "Content-Type: application/json" \
-d '{
    "$set": {
        "metadata.owner": "team-alpha",
        "public_url": "https://controller.example.com"
    }
}'

```

**Endpoint:** /vdag-controller/:controller\_id

**Method:** DELETE

**Description:**

Deletes the vDAG Controller document identified by the given controller\_id.

**Example curl Command:**

```
curl -X DELETE http://<server-url>/vdag-controller/controller-123
```

---

**Endpoint:** /vdag-controllers

**Method:** POST

**Description:**

Queries multiple vDAG Controller documents using a MongoDB-style filter object.

**Example curl Command:**

```
curl -X POST http://<server-url>/vdag-controllers \
  -H "Content-Type: application/json" \
  -d '{
    "cluster_id": "cluster-west-1",
    "metadata.owner": "team-alpha"
  }'
```

---

**Endpoint:** /vdag-controllers/by-vdag-uri/:vdag\_uri

**Method:** GET

**Description:**

Fetches all vDAG Controller documents associated with the given vdag\_uri.

**Example curl Command:**

```
curl -X GET http://<server-url>/vdag-controllers/by-vdag-uri/sample-vdag:1.0-stable
```

---